

Software-Defined 7 Networking and Advanced Network Control Programming

SDN Programming

Kai Gao

kaigao@scu.edu.cn

School of Cyber Science and Engineering
Sichuan University



Recap

Network Programming Language

Network programming language
in this course refer to
programming languages in the
SDN architecture

- Compile high-level intents to low-level device configurations (flow rule, BGP configuration, etc.)
- Maintain policy compliance under network events (topology change, traffic change, policy change, etc.)

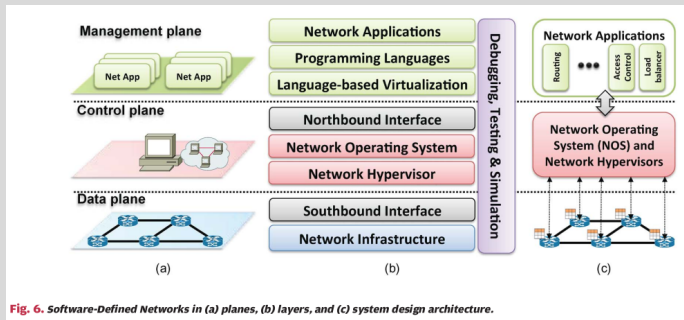


Fig. 6. Software-Defined Networks in (a) planes, (b) layers, and (c) system design architecture.

D. Kreutz et al. "Software-Defined Networking: A Comprehensive Survey". In: *Proceedings of the IEEE* 103.1 (Jan. 2015), pp. 14–76

Frenetic

Problems

- Error-prone interaction between individual components

Solutions:

Program 1: Simple switching

```
def switch_join(switch):
    repeater(switch)
def repeater(switch):
    pat1 = {in_port:1}
    pat2 = {in_port:2}
    install(switch,pat1,DEFAULT,None,[output(2)])
    install(switch,pat2,DEFAULT,None,[output(1)])
```

Program 2: Traffic monitoring

```
def monitor(switch):
    pat = {in_port:2,tp_src:80}
    install(switch,pat,DEFAULT,None,[])
    query_stats(switch,pat)
def stats_in(switch,xid,pattern,packets,bytes):
    print bytes
    sleep(30)
    query_stats(switch,pattern)
```

Frenetic

Problems

- Error-prone interaction between individual components
- Complex rule construction using low-level API

Solutions:

```
def repeater_monitor_noserver(switch):  
    pat1 = {in_port:1}  
    pat2 = {in_port:2}  
    pat2web = {in_port:2,tp_src:80}  
    pat2srv = {in_port:2,nw_dst:10.0.0.9,tp_src:80}  
    install(switch,pat1,DEFAULT,None,[output(2)])  
    install(switch,pat2srv,HIGH,None,[output(1)])  
    install(switch,pat2web,MEDIUM,None,[output(1)])  
    install(switch,pat2,DEFAULT,None,[output(1)])  
    query_stats(switch,pat2web)
```

Frenetic

Problems

- Error-prone interaction between individual components
- Complex rule construction using low-level API
- Race conditions when programming both control and data plane

Solutions:

Program

```
def repeater_monitor_hosts(switch):  
    pat = {in_port:1}  
    install(switch,pat,DEFAULT,None,[output(2)])  
def packet_in(switch,inport,packet):  
    if inport == 2:  
        mac = dstmac(packet)  
        pat = {in_port:2,dl_dst:mac}  
        install(switch,pat,DEFAULT,None,[output(1)])  
        query_stats(switch,pat)
```

Packets in arriving order:

in_port: 1, dl_dst: C5:85:2D:D6:B6:8B, tp_dst: 80

in_port: 1, dl_dst: C5:85:2D:D6:B6:8B, tp_dst: 8080

Flow table:

<i>match</i>	<i>priority</i>	<i>action</i>
in_port: 1, dl_dst: C5:85:2D:D6:B6:8B	DEFAULT	output: 2
in_port: 1, dl_dst: C5:85:2D:D6:B6:8B	DEFAULT	output: 2

Frenetic

Problems

- Error-prone interaction between individual components
- Complex rule construction using low-level API
- Race conditions when programming both control and data plane

Solutions:

- Network query language

```
Queries      q ::= Select(a) *  
              Where(fp) *  
              GroupBy([qh1, ..., qhn]) *  
              SplitWhen([qh1, ..., qhn]) *  
              Every(n) *  
              Limit(n)  
  
Aggregates  a ::= packets | sizes | counts  
Headers     qh ::= inport | srcmac | dstmac | ethtype |  
                vlan | srcip | dstip | protocol |  
                srcport | dstport | switch  
  
Patterns    fp ::= true_fp() | qh_fp(n) |  
              and_fp([fp1, ..., fpn]) |  
              or_fp([fp1, ..., fpn]) |  
              diff_fp(fp1, fp2) | not_fp(fp)
```

Figure 3. Frenetic query syntax

Frenetic

Problems

- Error-prone interaction between individual components
- Complex rule construction using low-level API
- Race conditions when programming both control and data plane

Solutions:

- Network query language
- Functional reactive network policy library

<i>Events</i>	
Seconds	$\in \text{int } E$
SwitchJoin	$\in \text{switch } E$
SwitchExit	$\in \text{switch } E$
PortChange	$\in (\text{switch} \times \text{int} \times \text{bool}) E$
Once	$\in \alpha \rightarrow \alpha E$
<i>Basic Event Functions</i>	
\gg	$\in \alpha E \rightarrow \alpha \beta EF \rightarrow \beta E$
Lift	$\in (a \rightarrow \beta) \rightarrow \alpha \beta EF$
\gg	$\in \alpha \beta EF \rightarrow \beta \gamma EF \rightarrow \alpha \gamma EF$
ApplyFst	$\in \alpha \beta EF \rightarrow (\alpha \times \gamma) (\beta \times \gamma) EF$
ApplySnd	$\in \alpha \beta EF \rightarrow (\gamma \times \alpha) (\gamma \times \beta) EF$
Merge	$\in (\alpha E \times \beta E) \rightarrow (\alpha \text{ option} \times \beta \text{ option}) E$
BlendLeft	$\in \alpha \times \alpha E \times \beta E \rightarrow (\alpha \times \beta) E$
BlendRight	$\in \beta \times \alpha E \times \beta E \rightarrow (\alpha \times \beta) E$
Accum	$\in (\gamma \times (\alpha \times \gamma \rightarrow \gamma)) \rightarrow \alpha \gamma EF$
Filter	$\in (\alpha \rightarrow \text{bool}) \rightarrow \alpha \alpha EF$
<i>Listeners</i>	
\gg	$\in \alpha E \rightarrow \alpha L \rightarrow \text{unit}$
Print	$\in \alpha L$
Register	$\in \text{policy } L$
Send	$\in (\text{switch} \times \text{packet} \times \text{action}) L$
<i>Rules and Policies</i>	
Rule	$\in \text{pattern} \times \text{action list} \rightarrow \text{rule}$
MakeForwardRules	$\in (\text{switch} \times \text{port} \times \text{packet}) \text{ policy } EF$
AddRules	$\in \text{policy policy } EF$

Figure 4. Selected Frenetic Operators.

Pyretic

Goal:

- Modular network policy

Solutions:

Monitor

srcip=5.6.7.8 → count

Route

dstip=10.0.0.1 → fwd(1)

dstip=10.0.0.2 → fwd(2)

Load-balance

srcip=0*,dstip=1.2.3.4 → dstip=10.0.0.1

srcip=1*,dstip=1.2.3.4 → dstip=10.0.0.2

Compiled Prioritized Rule Set for “Monitor | Route”

srcip=5.6.7.8,dstip=10.0.0.1 → count,fwd(1)

srcip=5.6.7.8,dstip=10.0.0.2 → count,fwd(2)

srcip=5.6.7.8 → count

dstip=10.0.0.1 → fwd(1)

dstip=10.0.0.2 → fwd(2)

Compiled Prioritized Rule Set for “Load-balance >> Route”

srcip=0*,dstip=1.2.3.4 → dstip=10.0.0.1,fwd(1)

srcip=1*,dstip=1.2.3.4 → dstip=10.0.0.2,fwd(2)

Goal:

- Modular network policy
- Programming on virtualized networks

Solutions:

Many-to-one Mapping:




Figure 2: Many physical switches to one virtual.

One-to-many Mapping:




Figure 3: One physical switch to many virtual.

Pyretic

Goal:

- Modular network policy
- Programming on virtualized networks

Solutions:

- Composition of network policies

Primitive Actions:

$A ::= \text{drop} \mid \text{passthrough} \mid \text{fwd}(\text{port}) \mid \text{flood} \mid \text{push}(h=v) \mid \text{pop}(h) \mid \text{move}(h1=h2)$

Predicates:

$P ::= \text{all_packets} \mid \text{no_packets} \mid \text{match}(h=v) \mid \text{ingress} \mid \text{egress} \mid P \ \& \ P \mid (P \mid P) \mid \sim P$

Query Policies:

$Q ::= \text{packets}(\text{limit}, [h]) \mid \text{counts}(\text{every}, [h])$

Policies:

$C ::= A \mid Q \mid P[C] \mid (C \mid C) \mid C \gg C \mid \text{if_}(P, C, C)$

Figure 4: Summary of static NetCore syntax.

Pyretic

Goal:

- Modular network policy
- Programming on virtualized networks

Solutions:

- Composition of network policies
- Network objects and policy transformations

```
def virtualize(ingress_policy,  
               egress_policy,  
               fabric_policy,  
               derived_policy):
```

Maple

Goal:

- Algorithmic policy
- Automatic & correct rule generation

Solution:

- Trace tree
- Barrier rule

```
def outgoing_policy(pkt):  
    if pkt.srcip in 192.168.0.0/24:  
        if pkt.srcport == 22:  
            use_path(s1, s3, s2)  
        else:  
            use_path(s1, s2)
```

Maple

Goal:

- Algorithmic policy
- Automatic & correct rule generation

Solution:

- Trace tree
- Barrier rule

```
def outgoing_policy(pkt):  
    if pkt.srcip in 192.168.0.0/24:  
        if pkt.srcport == 22:  
            use_path(s1, s3, s2)  
        else:  
            use_path(s1, s2)
```

<i>match</i>	<i>priority</i>	<i>action</i> (on s1)
srcip=192.168.0.0/24, srcport=22	MEDIUM	output(2)
srcip=192.168.0.0/24	DEFAULT	output(2)

Maple


Goal:

- Algorithmic policy
- Automatic & correct rule generation

Solution:

- Trace tree
- Barrier rule

```
def outgoing_policy(pkt):  
    if pkt.srcip in 192.168.0.0/24:  
        if pkt.srcport == 22:  
            use_path(s1, s3, s2)  
        else:  
            use_path(s1, s2)
```



Maple

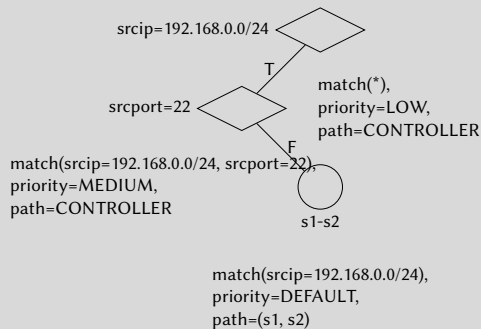
Goal:

- Algorithmic policy
- Automatic & correct rule generation

Solution:

- Trace tree
- Barrier rule

```
def outgoing_policy(pkt):  
    if pkt.srcip in 192.168.0.0/24:  
        if pkt.srcport == 22:  
            use_path(s1, s3, s2)  
        else:  
            use_path(s1, s2)
```



Magellan


Goal:

- *Proactively* build rules for algorithmic policies

Solutions:

- Data flow analysis
- DFG partition and cost analysis

```
def outgoing_policy(pkt):  
    if pkt.srcip in 192.168.0.0/24:  
        if pkt.srcport == 22:  
            use_path(s1, s3, s2)  
        else:  
            use_path(s1, s2)
```



Magellan


Goal:

- *Proactively* build rules for algorithmic policies

Solutions:

- Data flow analysis
- DFG partition and cost analysis

```
def policy(pkt):  
    srcvlan = mac2vlan[pkt.srcmac]  
    dstvlan = mac2vlan[pkt.dstmac]  
    use_path(vlan2path[srcvlan, dstvlan])
```



Magellan


Goal:

- *Proactively* build rules for algorithmic policies

Solutions:

- Data flow analysis
- DFG partition and cost analysis

```
def policy(pkt):  
    srcvlan = mac2vlan[pkt.srcmac]  
    dstvlan = mac2vlan[pkt.dstmac]  
    use_path(vlan2path[srcvlan, dstvlan])
```



Summary

In the coming lectures, we cover the following topics on network control programming

- SDN-based policy routing
- distributed enforcement of centralized policy
- programming stateful networks

Summary

In the coming lectures, we cover the following topics on network control programming

- SDN-based policy routing
- distributed enforcement of centralized policy
- programming stateful networks

In this lecture, you should

- learn network programming languages Merlin, Propane and SNAP

Summary

In the coming lectures, we cover the following topics on network control programming

- SDN-based policy routing
- distributed enforcement of centralized policy
- programming stateful networks

In this lecture, you should

- learn network programming languages Merlin, Propane and SNAP
- learn path constraints expressed in regular expressions

Summary

In the coming lectures, we cover the following topics on network control programming

- SDN-based policy routing
- distributed enforcement of centralized policy
- programming stateful networks

In this lecture, you should

- learn network programming languages Merlin, Propane and SNAP
- learn path constraints expressed in regular expressions
- learn how to use product graph to find policy compliant paths

Summary

In the coming lectures, we cover the following topics on network control programming

- SDN-based policy routing
- distributed enforcement of centralized policy
- programming stateful networks

In this lecture, you should

- learn network programming languages Merlin, Propane and SNAP
- learn path constraints expressed in regular expressions
- learn how to use product graph to find policy compliant paths
- learn how to use ILP to solve state placement problem

Merlin

Merlin: A Language for Provisioning Network Resources

Robert Soulé* Shrutarshi Basu† Parisa Jalili Marandi* Fernando Pedone*

Robert Kleinberg† Emin Gün Sirer† Nate Foster†

*University of Lugano †Cornell University

Robert Soulé et al. “Merlin: A Language for Provisioning Network Resources”. In: *Proceedings of the 10th ACM International on Conference on Emerging Networking Experiments and Technologies*. CoNEXT '14. New York, NY, USA: ACM, 2014, pp. 213–226. URL: <http://doi.acm.org/10.1145/2674005.2674989>

Goal

Support **resource provisioning** in SDN programming languages

There are two types of resource constraints:

- **Guarantee:** the minimum bandwidth the traffic can reach
- **Constraint:** the maximum bandwidth the traffic can reach

Example: Assume there are two flows f_1 and f_2 traversing a link with 5 Mbps

- rate of f_1 : r_1 , rate of f_2 : r_2
- bandwidth of f_1 : b_1 , bandwidth of f_2 : b_2
- guarantee of f_1 : 2 Mbps
- constraint of f_1 : 4 Mbps

Goal

Support **resource provisioning** in SDN programming languages


There are two types of resource constraints:

- **Guarantee:** the minimum bandwidth the traffic can reach
- **Constraint:** the maximum bandwidth the traffic can reach

Example: Assume there are two flows f_1 and f_2 traversing a link with 5 Mbps

- rate of f_1 : r_1 , rate of f_2 : r_2
- bandwidth of f_1 : b_1 , bandwidth of f_2 : b_2
- guarantee of f_1 : 2 Mbps
- constraint of f_1 : 4 Mbps

Feasible region of rate:



Goal

Support **resource provisioning** in SDN programming languages


There are two types of resource constraints:

- **Guarantee:** the minimum bandwidth the traffic can reach
- **Constraint:** the maximum bandwidth the traffic can reach

Example: Assume there are two flows f_1 and f_2 traversing a link with 5 Mbps

- rate of f_1 : r_1 , rate of f_2 : r_2
- bandwidth of f_1 : b_1 , bandwidth of f_2 : b_2
- guarantee of f_1 : 2 Mbps
- constraint of f_1 : 4 Mbps

Feasible region of bandwidth:



Merlin Language

Merlin allows programmers to specify policies with resource requirements

Robert Soulé et al. "Merlin: A Language for Provisioning Network Resources". In: *Proceedings of the 10th ACM International on Conference on Emerging Networking Experiments and Technologies*. CoNEXT '14. New York, NY, USA: ACM, 2014, pp. 213–226. URL: <http://doi.acm.org/10.1145/2674005.2674989>

$loc \in$	<i>Locations</i>	
$t \in$	<i>Packet-processing functions</i>	
$h \in$	<i>Packet headers</i>	
$f \in$	<i>Header fields</i>	
$v \in$	<i>Header field values</i>	
$id \in$	<i>Identifiers</i>	
$n \in$	\mathbb{N}	
$pol ::=$	$[s_1; \dots; s_n], \phi$	Policies
$s ::=$	$id : p \rightarrow r$	Statements
$\phi ::=$	$\max(e, n) \mid \min(e, n)$ $\mid \phi_1 \text{ and } \phi_2 \mid \phi_1 \text{ or } \phi_2 \mid ! \phi_1$	Presburger Formulas
$e ::=$	$n \mid id \mid e + e$	Bandwidth Terms
$a ::=$	$\cdot \mid c \mid a a \mid a \mid a \mid a^* \mid ! a$	Path Expression
$p ::=$	$p_1 \text{ and } p_2 \mid p_1 \text{ or } p_2 \mid ! p_1$ $\mid h.f = v \mid \text{true} \mid \text{false}$	Predicates
$c ::=$	$loc \mid t$	Path Element

Figure 1: Merlin abstract syntax.

Example

```
[ x : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 20) -> .* dpi .* ;
  y : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 21) -> .* ;
  z : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 80) -> .* dpi .* nat .* ],
max(x + y, 50MB/s) and min(z, 100MB/s)
```

- This policy defines 3 flows: x , y and z
- Total bandwidth of flows x and y must not exceed 50 MB/s
- Bandwidth of flow z has a guarantee of 100 MB/s

A Closer Look at Flow Specification

```
x : (ip.src = 192.168.1.1 and
     ip.dst = 192.168.1.2 and
     tcp.dst = 20) -> .* dpi .* ;
```

- x: identifier
- ip.src = 192.168.1.1 and ip.dst = 192.168.1.2 and tcp.dst = 20: predicate
- .* dpi .*: path expression

$loc \in \text{Locations}$

$t \in \text{Packet-processing functions}$

$h \in \text{Packet headers}$

$f \in \text{Header fields}$

$v \in \text{Header field values}$

$id \in \text{Identifiers}$

$n \in \mathbb{N}$

$pol ::= [s_1; \dots; s_n], \phi$

Policies

$s ::= id : p \rightarrow r$

Statements

$\phi ::= \max(e, n) \mid \min(e, n) \mid \phi_1 \text{ and } \phi_2 \mid \phi_1 \text{ or } \phi_2 \mid !\phi_1$

Presburger Formulas

$e ::= n \mid id \mid e + e$

Bandwidth Terms

$a ::= . \mid c \mid aa \mid a \mid a^* \mid !a$

Path Expression

$p ::= p_1 \text{ and } p_2 \mid p_1 \text{ or } p_2 \mid !p_1 \mid h.f = v \mid \text{true} \mid \text{false}$

Predicates

$c ::= loc \mid t$

Path Element

Figure 1: Merlin abstract syntax.

Path Expression

Merlin uses **path expression** to specify the **traversal constraint**

Merlin path expression is a **regular language** and can be represented as a finite state automaton

Example


Path Expression

Merlin uses **path expression** to specify the traversal constraint

Merlin path expression is a **regular language** and can be represented as a finite state automaton

Example

- `.* dpi .*`





Path Expression

Merlin uses **path expression** to specify the **traversal constraint**


Merlin path expression is a **regular language** and can be represented as a finite state automaton

Example

- `.* dpi .*`
- `.* dpi .* nat .*`



Logical Topology



Policy:


`h1 .* dpi .* nat .* h2`

Locations:

- dpi: h_1, h_2 , or m_1
- nat: m_1

Figure 2: Logical topology for the example policy. The thick red path illustrates a solution.

Logical Topology



Policy:

h1 .* dpi .* nat .* h2

Locations:

• dpi: h1, h2, or m1


• nat: m1

Path:

h1

Figure 2: Logical topology for the example policy. The thick red path illustrates a solution.

Logical Topology



Policy:

h1 .* dpi .* nat .* h2

Locations:

• dpi: h1, h2, or m1


• nat: m1

Path:

h1

Figure 2: Logical topology for the example policy. The thick red path illustrates a solution.

Logical Topology



Policy:

h1 .* **dpi** .* nat .* h2

Locations:

• dpi: h1, h2, or m1


• nat: m1

Path:

h1

Figure 2: Logical topology for the example policy. The thick red path illustrates a solution.

Logical Topology



Policy:

`h1 .* dpi .* nat .* h2`

Locations:

- dpi: h_1, h_2 , or m_1


- nat: m_1

Path:

`h1 s1`

Figure 2: Logical topology for the example policy. The thick red path illustrates a solution.

Logical Topology



Policy:

`h1 .* dpi .* nat .* h2`

Locations:

- dpi: h_1, h_2 , or m_1


- nat: m_1

Path:

`h1 s1 m1`

Figure 2: Logical topology for the example policy. The thick red path illustrates a solution.

Logical Topology



Policy:

`h1 .* dpi .* nat .* h2`

Locations:

- dpi: h_1, h_2 , or m_1


- nat: m_1

Path:

`h1 s1 m1 s1`

Figure 2: Logical topology for the example policy. The thick red path illustrates a solution.

Logical Topology



Policy:

`h1 .* dpi .* nat .* h2`

Locations:

- dpi: h1, h2, or m1

- nat: m1


Path:

`h1 s1 m1 s1 h2`

Figure 2: Logical topology for the example policy. The thick red path illustrates a solution.

Policy Localization


Bandwidth terms such as $\max(x + y, 50 \text{ MB/s})$ cannot be realized on hardware if they take different paths



Policy Localization

Bandwidth terms such as $\max(x + y, 50 \text{ MB/s})$ cannot be realized on hardware if they take different paths


- $a \rightarrow b \rightarrow d$, $a \rightarrow c \rightarrow d$: can be enforced at a or d



Policy Localization

Bandwidth terms such as $\max(x + y, 50 \text{ MB/s})$ cannot be realized on hardware if they take different paths

- $a \rightarrow b \rightarrow d$, $a \rightarrow c \rightarrow d$: can be enforced at a or d
- $a \rightarrow c$, $b \rightarrow d$: cannot be enforced together



Policy Localization

Merlin transforms original policy into **localized policy**

```
[ x : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 20) -> .* dpi .* ;
  y : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 21) -> .* ;
  z : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 80) -> .* dpi .* nat .* ],
  max(x + y, 50MB/s) and min(z, 100MB/s)
```

```
[ x : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 20) -> .* dpi .* ;
  y : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 21) -> .* ;
  z : (ip.src = 192.168.1.1 and
      ip.dst = 192.168.1.2 and
      tcp.dst = 80) -> .* dpi .* nat .* ],
  max(x, 25MB/s) and max(y, 25MB/s) and min(z,
```

Global constraint $\max(x + y, 50\text{MB/s})$ is transformed to localized constraint $\max(x, 25\text{MB/s})$ and $\max(y, 25\text{MB/s})$

Policy Localization

- Constraints on multiple flows are transformed to constraints on **single flows**
- Localized constraints **must** still satisfy the global constraints

Example:

$\min(x + y, 100 \text{ MB/s})$ and $\min(y + z, 60 \text{ MB/s})$ and $\max(y, 40 \text{ MB/s})$

✗ $\min(x, 50\text{MB/s})$ and $\min(y, 50\text{MB/s})$ and $\min(y, 30\text{MB/s})$ and $\min(z, 30\text{MB/s})$

Policy Localization

- Constraints on multiple flows are transformed to constraints on **single flows**
- Localized constraints **must** still satisfy the global constraints

Example:

$\min(x + y, 100 \text{ MB/s})$ and $\min(y + z, 60 \text{ MB/s})$ and $\max(y, 40 \text{ MB/s})$

✗ $\min(x, 50\text{MB/s})$ and $\min(y, 50\text{MB/s})$ and $\min(y, 30\text{MB/s})$ and $\min(z, 30\text{MB/s})$

✓ $\min(x, 70\text{MB/s})$ and $\min(y, 30\text{MB/s})$ and $\min(z, 30\text{MB/s})$

Policy Localization

- Constraints on multiple flows are transformed to constraints on **single flows**
- Localized constraints **must** still satisfy the global constraints

Example:

$\min(x + y, 100 \text{ MB/s})$ and $\min(y + z, 60 \text{ MB/s})$ and $\max(y, 40 \text{ MB/s})$

✗ $\min(x, 50\text{MB/s})$ and $\min(y, 50\text{MB/s})$ and $\min(y, 30\text{MB/s})$ and $\min(z, 30\text{MB/s})$

✓ $\min(x, 70\text{MB/s})$ and $\min(y, 30\text{MB/s})$ and $\min(z, 30\text{MB/s})$

✓ $\min(x, 60\text{MB/s})$ and $\min(y, 40\text{MB/s})$ and $\min(z, 20\text{MB/s})$

Find Solution using (Mixed) Integer Linear Programming

Decision variables:

- $x_e \in \{0, 1\}$ - whether an edge in the **logical topology** will be selected
- $r_{uv} \in [0, 1]$ - link utilization on link (u, v)
- $r_{max} \in [0, 1]$ - maximum link utilization
- $R_{max} \in [0, C]$ - maximum link rate

Flow conservation constraints:

$$\forall v \in \mathcal{G}, \sum_{e \in \delta^+(v)} x_e - \sum_{e \in \delta^-(v)} x_e = \begin{cases} 1 & \text{if } v = s_i \\ -1 & \text{if } v = t_i \\ 0 & \text{otherwise} \end{cases}$$

Resource constraints:

$$\forall (u, v), r_{uv} c_{uv} = \sum_i \sum_{e \in E_i(u, v)} r_{min}^i x_e \quad (\text{definition of link utilization})$$

$$\forall (u, v), r_{max} \geq r_{uv} \quad (\text{definition of } r_{max})$$

$$\forall (u, v), R_{max} \geq r_{uv} c_{uv} \quad (\text{definition of } R_{max})$$

$$r_{max} \leq 1 \quad (\text{link capacity})$$

Summary

- Merlin is an SDN programming language that enables programmers to specify resource constraints
- Merlin uses logical topology (cross product of topology and policy automaton) to find paths that satisfy waypoint traversal constraints
- Merlin transforms global constraints into localized constraints
- Merlin uses MILP to find feasible solution that satisfy the specified constraints

Propane

Don't Mind the Gap: Bridging Network-wide Objectives and Device-level Configurations

Ryan Beckett
Princeton

Ratul Mahajan
Microsoft

Todd Millstein
UCLA


Jitendra Padhye
Microsoft

David Walker
Princeton

Ryan Beckett et al. "Don't Mind the Gap: Bridging Network-wide Objectives and Device-level Configurations". In: *Proceedings of the 2016 ACM SIGCOMM Conference*. SIGCOMM '16. New York, NY, USA: ACM, 2016, pp. 328–341. URL: <http://doi.acm.org/10.1145/2934872.2934909>

Problems

Complexity of using Decentralized Routing




Policy

- P1. Prefer Cust > Peer > Prov
- P2. Disallow traffic between Prov and Peer
- P3. For Cust, prefer R1 > R2
- P4. Cust must be on path for its prefixes
- P5. Cust must not be a transit to Prov

Figure 1: Creating router-level policies is difficult.

Problems

Complexity when Failures Happen



Policy

P1. Left cluster has global services with PG* prefixes, which should be announced externally as an aggregate PG

P2. Right cluster has local services with PL* prefixes, which should not be announced externally

Figure 2: Policy-compliance under failures is difficult.

Propane Example 1

```
define Prefs = exit(R1 » R2 » Peer » Prov)
```

(P1 and P3)

```
define Routing =  
  {PCust => Prefs & end(Cust)  
  true  => Prefs }
```


(P4, P1 and P3)

```
define transit(X,Y) = enter(X|Y) & exit(X|Y)  
define cust-transit(X,Y) = later(X) & later(Y)
```

```
define NoTrans =  
  {true => !transit(Peer,Prov) &  
   !cust-transit(Cust,Prov) }
```

(P2, P5)

Final policy: Routing & NoTrans



Policy

- P1. Prefer Cust > Peer > Prov
- P2. Disallow traffic between Prov and Peer
- P3. For Cust, prefer R1 > R2
- P4. Cust must be on path for its prefixes
- P5. Cust must not be a transit to Prov


Figure 1: Creating router-level policies is difficult.

Propane Example 2

```
define Ownership =  
  {PG1 => end(A)  
   PG2 => end(B)  
   PL1 => end(E)  
   PL2 => end(F)  
   true => end(out) }
```

```
define Locality =  
  {PL1 | PL2 => only(in) }
```

Final policy: Ownership & Locality



Policy

P1. Left cluster has global services with PG* prefixes, which should be announced externally as an aggregate PG

P2. Right cluster has local services with PL* prefixes, which should not be announced externally

Figure 2: Policy-compliance under failures is difficult.

Propane Language

Syntax	Propane Expansions
$pol ::= p_1, \dots, p_n$	any = $out^* \cdot in^+ \cdot out^*$
$p ::= t \Rightarrow r_1 \gg \dots \gg r_m \mid cc$	drop = \emptyset
$x ::= d.d.d.d/d$	internal = in^+
$t ::= true$	only (X) = $any \cap X^*$
$\quad \mid !t$	never (X) = $any \cap (!X)^*$
$\quad \mid t_1 \mid t_2$	through (X) = $out^* \cdot in^* \cdot X \cdot in^* \cdot out^*$
$\quad \mid t_1 \& t_2$	later (X) = $out^* \cdot (X \cap out) \cdot out^* \cdot in^+ \cdot out^*$
$\quad \mid prefix = x$	before (X) = $out^* \cdot in^+ \cdot out^* \cdot (X \cap out) \cdot out^*$
$\quad \mid comm = d$	end (X) = $any \cap (\Sigma^* \cdot X)$
$r ::= l$	start (X) = $any \cap (X \cdot \Sigma^*)$
$\quad \mid \emptyset$	exit (X) = $(out^* \cdot in^* \cdot (X \cap in) \cdot out \cdot out^*) \cup$ $(out^* \cdot in^+ \cdot (X \cap out) \cdot out^*)$
$\quad \mid in$	enter (X) = $(out^* \cdot out \cdot (X \cap in) \cdot in^* \cdot out^*) \cup$ $(out^* \cdot (X \cap out) \cdot in^+ \cdot out^*)$
$\quad \mid out$	link (X, Y) = $any \cap (\Sigma^* \cdot X \cdot Y \cdot \Sigma^*)$
$\quad \mid r_1 \cup r_2$	path (\vec{X}) = $any \cap (\Sigma^* \cdot X_1 \dots X_n \cdot \Sigma^*)$
$\quad \mid r_1 \cap r_2$	novalley (\vec{X}) = $any \cap$ $!path(X_2, X_1, X_2) \cap \dots \cap$ $!path(X_n, X_{n-1}, X_n)$
$\quad \mid r_1 \cdot r_2$	
$\quad \mid !r$	
$\quad \mid r^*$	
$ln ::= r_1 \rightarrow r_2$	
$cc ::= agg(x, ln) \mid tag(d, t, ln)$	

Figure 4: Regular Intermediate Representation (RIR) syntax (left), and Propane language expansions (right).

Propane Compilation

We focus on the translations from FE to RIR and from RIR to PGIR




Figure 3: Compilation pipeline stages for Propane.

Regular Intermediate Representation (RIR)

- Each rule has the format `prefix => path`
- `prefix` is a set of IPv4 prefixes
- `path` denotes a set of path that satisfies the expression
- Policy composition is similar to OpenFlow policy composition
- After **rule-based rewriting**, the policy can be rewritten as one policy where each rule specifies **the path for a single prefix**

Example:

```
define Ownership =  
  {PG1 => end(A)  
   PG2 => end(B)  
   PL1 => end(E)  
   PL2 => end(F)  
   true => end(out) }
```

```
define Locality =  
  {PL1 | PL2 => only(in) }
```

Ownership & Locality:

```
PG1 => end(A)  
PG2 => end(B)  
PL1 => only(in)  $\cap$  end(E)  
PL2 => only(in)  $\cap$  end(F)  
true => exit(out)
```


Path Expression to Automaton

Propane path expression is a regular language and can be represented as a finite state automaton



Example:

- $W \cdot A \cdot C \cdot D \cdot \text{out}$
- $W \cdot B \cdot \text{in}^+ \cdot \text{out}$

Topology




Policy Automata




Product Graph IR


Topology



Policy Automata



Product Graph IR




Product Graph IR

Topology




- Each state is a triplet $N, S1, S2$ where N is the node, $S1$ is the state in automata 1, $S2$ is the state in automata 2

Policy Automata




Product Graph IR




Product Graph IR

Topology

- Each state is a triplet $N, S1, S2$ where N is the node, $S1$ is the state in automata 1, $S2$ is the state in automata 2
- Solid line represents a valid transition




Product Graph IR




Product Graph IR

Topology

- Each state is a triplet $N, S1, S2$ where N is the node, $S1$ is the state in automata 1, $S2$ is the state in automata 2
- Solid line represents a valid transition
- Dashed line represents an invalid transition



Product Graph IR




Product Graph IR


Topology



- Each state is a triplet $N, S1, S2$ where N is the node, $S1$ is the state in automata 1, $S2$ is the state in automata 2
- Solid line represents a valid transition
- Dashed line represents an invalid transition
- Red nodes are eliminated because there are no solid paths using these nodes



Product Graph IR

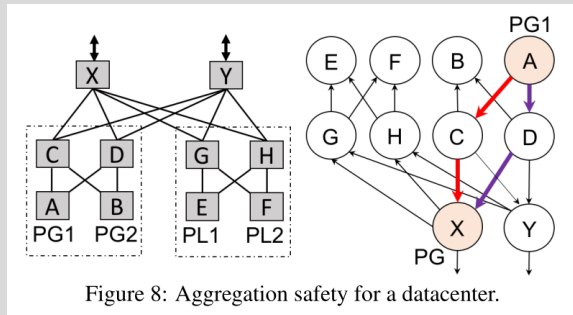


What's Next

After obtaining the PGIR, Propane conducts **preference inference** with analysis on:

- **failure-safety**: whether the policy can be realized under network failures
- **aggregation-safety**: whether sub-prefixes can be aggregated into a shorter prefix

Then PGIR is translated into **abstract BGP configuration**



Summary

- Propane is a programming language that allows programmers to specify **centralized prefix routing requirements**
- Propane policies are organized as well-formed **regular intermediate representation** (RIR)
- RIR is compiled with the topology to generate **product graph intermediate representation** (PGIR)
- Propane infers the local preferences with safety analysis, and translates PGIR into abstract BGP

SNAP

SNAP: Stateful Network-Wide Abstractions for Packet Processing

Mina Tahmasbi Arashloo¹, Yaron Koral¹, Michael Greenberg², Jennifer Rexford¹, and David Walker¹

¹Princeton University, ²Pomona College

Mina Tahmasbi Arashloo et al. “SNAP: Stateful Network-Wide Abstractions for Packet Processing”. In: *Proceedings of the 2016 ACM SIGCOMM Conference. SIGCOMM '16*. Florianopolis, Brazil: Association for Computing Machinery, 2016, pp. 29–43. URL: <https://doi.org/10.1145/2934872.2934892>

Stateful Network Policies

Example: DNS tunnel detection

- Keep a counter for **resolved but not used** DNS requests **for each host**
- If a counter exceeds a threshold, block the traffic **from the host**

DNS-tunnel-detect

```
1  if dstip = 10.0.6.0/24 & srcport = 53 then
2    orphan[dstip][dns.rdata] <- True;
3    susp-client[dstip]++;
4    if susp-client[dstip] = threshold then
5      blacklist[dstip] <- True
6    else id
7  else
8    if srcip = 10.0.6.0/24 & orphan[srcip][dstip]
9      then orphan[srcip][dstip] <- False;
10     susp-client[srcip]--
11  else id
```

Figure 1: SNAP implementation of DNS-tunnel-detect.

Stateful Network Policies

Example: DNS tunnel detection

- Keep a counter for **resolved but not used** DNS requests **for each host**
- If a counter exceeds a threshold, block the traffic **from the host**

DNS-tunnel-detect

```
1  if dstip = 10.0.6.0/24 & srcport = 53 then
2    orphan[dstip][dns.rdata] <- True;
3    susp-client[dstip]++;
4    if susp-client[dstip] = threshold then
5      blacklist[dstip] <- True
6    else id
7  else
8    if srcip = 10.0.6.0/24 & orphan[srcip][dstip]
9      then orphan[srcip][dstip] <- False;
10     susp-client[srcip]--
11  else id
```

Figure 1: SNAP implementation of DNS-tunnel-detect.

- `orphan[dstip][dns.rdata]`: for each host (dstip), the request to `dns.rdata` is resolved but not used
- `susp-client[ip]`: the counter for host with ip

SNAP: Syntax

The syntax of SNAP:

- is based on the syntax of Pyretic
- with extensions to support stateful policies
- has no support for traversal order (one-big-switch abstraction)

$e \in \text{Expr}$	$::=$	$v \mid f \mid \vec{e}$	
$x, y \in \text{Pred}$	$::=$	id	Identity
		$drop$	Drop
		$f = v$	Test
		$\neg x$	Negation
		$x \mid y$	Disjunction
		$y \& x$	Conjunction
		$s[e] = e$	State Test
$p, q \in \text{Pol}$	$::=$	x	Filter
		$f \leftarrow v$	Modification
		$p + q$	Parallel comp.
		$p; q$	Sequential comp.
		$s[e] \leftarrow e$	State Modification
		$s[e]++$	Increment value
		$s[e]--$	Decrement value
		$\text{if } a \text{ then } p \text{ else } q$	Conditional
		$\text{atomic}(p)$	Atomic blocks

Figure 4: SNAP's syntax. Highlighted items are not in NetCore.

SNAP: Compilation

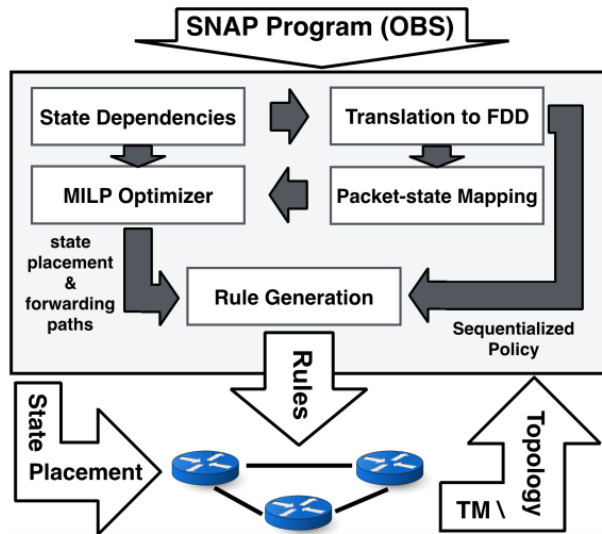


Figure 5: Overview of the compiler phases.

Main process:

- Translate SNAP program into xFDD
- Use xFDD to derive placement constraints for state variables
- Build MILP problem to solve the state placement problem
- Apply state placement and routing

Extended Forwarding Decision Diagram (xFDD) is similar to *binary decision diagram*

It captures the data dependencies in a network policy

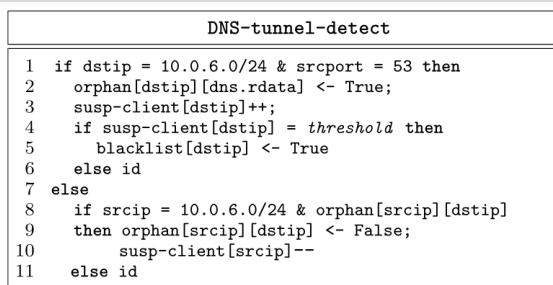


Figure 1: SNAP implementation of DNS-tunnel-detect.

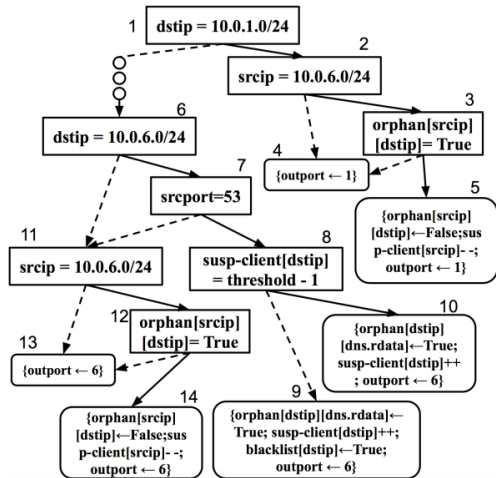


Figure 3: The equivalent xFDD for
DNS-tunnel-detect; assign-egress

xFDD Translation

SNAP uses **syntax-driven translation**

d	$::=$	$t ? d_1 : d_2 \mid \{as_1, \dots, as_n\}$	xFDDs
t	$::=$	$f = v \mid f_1 = f_2 \mid s[e_1] = e_2$	tests
as	$::=$	$a \mid a; a$	action sequences
a	$::=$	$id \mid drop \mid f \leftarrow v \mid s[e_1] \leftarrow e_2$ $\mid s[e_1]++ \mid s[e_1]--$	actions

$$\begin{aligned} \text{TO-XFDD}(a) &= \{a\} \\ \text{TO-XFDD}(f = v) &= f = v ? \{id\} : \{drop\} \\ \text{TO-XFDD}(\neg x) &= \ominus \text{TO-XFDD}(x) \\ \text{TO-XFDD}(s[e_1] = e_2) &= s[e_1] = e_2 ? \{id\} : \{drop\} \\ \text{TO-XFDD}(\text{atomic}(p)) &= \text{TO-XFDD}(p) \\ \text{TO-XFDD}(p + q) &= \text{TO-XFDD}(p) \oplus \text{TO-XFDD}(q) \\ \text{TO-XFDD}(p; q) &= \text{TO-XFDD}(p) \odot \text{TO-XFDD}(q) \\ \text{TO-XFDD}(\text{if } x \text{ then } p \text{ else } q) &= (\text{TO-XFDD}(x) \odot \text{TO-XFDD}(p)) \\ &\quad \oplus (\ominus \text{TO-XFDD}(x) \odot \text{TO-XFDD}(q)) \end{aligned}$$

Figure 6: xFDD syntax and translation.

NOTE: Policy compositions are translated to xFDD compositions

xFDD Composition

$\{as_{11}, \dots, as_{1n}\} \oplus \{as_{21}, \dots, as_{2m}\} = \{as_{11}, \dots, as_{1n}\} \cup \{as_{21}, \dots, as_{2m}\}$ $(t \text{ ? } d_1 : d_2) \oplus \{as_1, \dots, as_n\} = (t \text{ ? } d_1 \oplus \{as_1, \dots, as_n\} : d_2 \oplus \{as_1, \dots, as_n\})$		$\ominus\{id\} = \{drop\}$ $\ominus\{drop\} = \{id\}$ $\ominus(t?d_1 : d_2) = (t? \ominus d_1 : \ominus d_2)$
$(t_1 \text{ ? } d_{11} : d_{12}) \oplus (t_2 \text{ ? } d_{21} : d_{22}) = \begin{cases} (t_1 \text{ ? } d_{11} \oplus d_{21} : d_{12} \oplus d_{22}) & t_1 = t_2 \\ (t_1 \text{ ? } d_{11} \oplus (t_2 \text{ ? } d_{21} : d_{22}) : d_{12} \oplus (t_2 \text{ ? } d_{21} : d_{22})) & t_1 \sqsubset t_2 \\ (t_2 \text{ ? } d_{21} \oplus (t_1 \text{ ? } d_{11} : d_{12}) : d_{22} \oplus (t_1 \text{ ? } d_{11} : d_{12})) & t_2 \sqsubset t_1 \end{cases}$		
$as \odot \{as_1, \dots, as_n\} = \{as \odot as_1, \dots, as \odot as_n\}$ $as \odot (t \text{ ? } d_1 : d_2) = \text{(see explanations in §4.2)}$ $\{as_1, \dots, as_n\} \odot d = (as_1 \odot d) \oplus \dots \oplus (as_n \odot d)$ $(t \text{ ? } d_1 : d_2) \odot d = (d_1 \odot d) _t \oplus (d_2 \odot d) _{\sim t}$		$\{as_1, \dots, as_n\} _t = (t \text{ ? } \{as_1, \dots, as_n\} : \{drop\})$ $(t_1 \text{ ? } d_1 : d_2) _{t_2} = \begin{cases} (t_1 \text{ ? } d_1 : \{drop\}) & t_1 = t_2 \\ (t_2 \text{ ? } (t_1 \text{ ? } d_1 : d_2) : \{drop\}) & t_2 \sqsubset t_1 \\ (t_1 \text{ ? } d_1 _{t_2} : d_2 _{t_2}) & t_1 \sqsubset t_2 \end{cases}$

Figure 7: Definitions of xFDD composition operators.

State Placement and Routing Problem

SNAP places the program for **one-big-switch** in the network

DNS-tunnel-detect

```
1  if dstip = 10.0.6.0/24 & srcport = 53 then
2    orphan[dstip][dns.rdata] <- True;
3    susp-client[dstip]++;
4    if susp-client[dstip] = threshold then
5      blacklist[dstip] <- True
6    else id
7  else
8    if srcip = 10.0.6.0/24 & orphan[srcip][dstip]
9      then orphan[srcip][dstip] <- False;
10     susp-client[srcip]--
11  else id
```

Figure 1: SNAP implementation of DNS-tunnel-detect.




Figure 2: Topology for the running example.

This is achieved by constructing an MILP problem

Problem Formulation

Variables and Constants

Decision variables

- $R_{uvij} \in \{0, 1\}$: flow from u to v that traverses link (i, j)
- $P_{sn} \in \{0, 1\}$: state variable s is placed on node n
- $P_{suvij} \in \{0, 1\}$: flow from u to v has traversed state variable s before traversing link (i, j)

Constants

- d_{uv} : traffic from u to v
- c_{ij} : link capacity for link (i, j)
- S_{uv} : state variables that flow from u to v depends on
- state dependencies: $(s, t) \in tied$ - state variables s and t must be put on the same device
- state dependencies: $(s, t) \in dep$ - state variables s must be visited before t

Problem Formulation

Routing Constraints

$$\sum_j R_{uvuj} = 1 \quad \forall u, v \text{ (source)}$$

$$\sum_i R_{uviv} = 1 \quad \forall u, v \text{ (sink)}$$

$$\sum_{u,v} R_{uvij} d_{uv} \leq c_{ij} \quad \forall i, j \text{ (link capacity)}$$

$$\sum_i R_{uvin} = \sum_j R_{uvnj} \quad \forall u, v, n \text{ (flow conservation)}$$

$$\sum_i R_{uvin} \leq 1 \quad \forall u, v, n \text{ (single path)}$$

Problem Formulation

State Constraints

State constraints

$$\sum_n P_{sn} = 1 \quad \forall s \text{ (visited once)}$$

$$\sum_i R_{uvin} \geq P_{sn} \quad \forall u, v, \forall s \in S_{uv} \text{ (only placed on path)}$$

$$P_{sn} = P_{tn} \quad \forall (s, t) \in \textit{tied} \text{ (state placement constraint)}$$

$$P_{suvij} \leq R_{uvij} \quad \forall u, v, s \in S_{uv} \text{ (only traversed on path)}$$

$$P_{sn} + \sum_i P_{suvin} = \sum_j P_{suvij} \quad \forall u, v, s \in S_{uv} \text{ (state traversal)}$$

$$P_{sv} + \sum_i P_{suviv} = 1 \quad \forall u, v, s \in S_{uv} \text{ (state sink)}$$

$$P_{sn} + \sum_i P_{suvin} \geq P_{tn} \quad \forall u, v, (s, t) \in \textit{dep} \text{ (state dependency)}$$

Summary

- SNAP enables stateful network policies
- SNAP translates one-big-switch policies into xFDD to analyze the state dependencies
- The routing and state placement are modeled as an MILP problem

The End

Summary

In the coming lectures, we cover the following topics on network control programming

- SDN-based policy routing
- distributed enforcement of centralized policy
- programming stateful networks

In this lecture, you should

- learn network programming languages Merlin, Propane and SNAP
- learn path constraints expressed in regular expressions
- learn how to use product graph to find policy compliant paths
- learn how to use ILP to solve state placement problem

Thanks!

kaigao@scu.edu.cn

References I

- [1] Mina Tahmasbi Arashloo et al. “SNAP: Stateful Network-Wide Abstractions for Packet Processing”. In: *Proceedings of the 2016 ACM SIGCOMM Conference. SIGCOMM '16*. Florianopolis, Brazil: Association for Computing Machinery, 2016, pp. 29–43. URL: <https://doi.org/10.1145/2934872.2934892>.
- [2] Ryan Beckett et al. “Don’t Mind the Gap: Bridging Network-wide Objectives and Device-level Configurations”. In: *Proceedings of the 2016 ACM SIGCOMM Conference. SIGCOMM '16*. New York, NY, USA: ACM, 2016, pp. 328–341. URL: <http://doi.acm.org/10.1145/2934872.2934909>.
- [3] D. Kreutz et al. “Software-Defined Networking: A Comprehensive Survey”. In: *Proceedings of the IEEE* 103.1 (Jan. 2015), pp. 14–76.
- [4] Robert Soulé et al. “Merlin: A Language for Provisioning Network Resources”. In: *Proceedings of the 10th ACM International on Conference on Emerging Networking Experiments and Technologies. CoNEXT '14*. New York, NY, USA: ACM, 2014, pp. 213–226. URL: <http://doi.acm.org/10.1145/2674005.2674989>.